



095946- ADVANCED ALGORITHMS AND PARALLEL PROGRAMMING

Fabrizio Ferrandi

a.a. 2021-2022



-
- Randomized data structures
 - Random Treaps
 - Skip Lists



Dictionaries

- ❑ A dictionary is a collection of elements each of which has a **unique search key**
 - Uniqueness criteria may be relaxed (multiset)
 - (i.e. do not force uniqueness)
- ❑ Keep track of current members, with periodic insertions and deletions into the set
- ❑ Examples
 - Membership in a club, course records
 - Symbol table (contains duplicates)
 - Language dictionary (WordSmith, Webster, WordNet)
- ❑ Similar to database



Course Records

Dictionary

Member
Record →

key	student name	hw1	...
123	Stan Smith	49	...
124	Sue Margolin	56	...
125	Billie King	34	...
⋮			
167	Roy Miller	39	...

Satelite data



The dictionary problem

- **Given:** Universe $(U, <)$ of keys with a total order
- **Goal:** Maintain set $S \subseteq U$ under the following operations
 - **Search(x, S):** Is $x \in S$?
 - **Insert(x, S):** Insert x into S if not already in S .
 - **Delete(x, S):** Delete x from S .



Extended set of operations

possible problems to do with dictionaries

- **Minimum(S)**: Return smallest key
- **Maximum(S)**: Return largest key
- **List(S)**: Output elements of S in increasing order by key
- **Union(S_1, S_2)**: Merge S_1 and S_2
Condition: $\forall x_1 \in S_1, x_2 \in S_2: x_1 < x_2$
- **Split(S, x, S_1, S_2)**: Split S into S_1 and S_2 .
 $\forall x_1 \in S_1, x_2 \in S_2: x_1 \leq x$ and $x < x_2$



How to Implement a Dictionary?

- Sequences / Arrays
 - ordered
 - unordered
- Binary Search Trees
- AVL tree
- Splay tree

- Random Treaps
- Skip lists



Recall Arrays ...

- Unordered array

- *unordered sequence*



- searching and removing takes $O(n)$
- inserting takes $O(1)$ time
- applications to log files (frequent insertions, rare searches and removals)



More Arrays

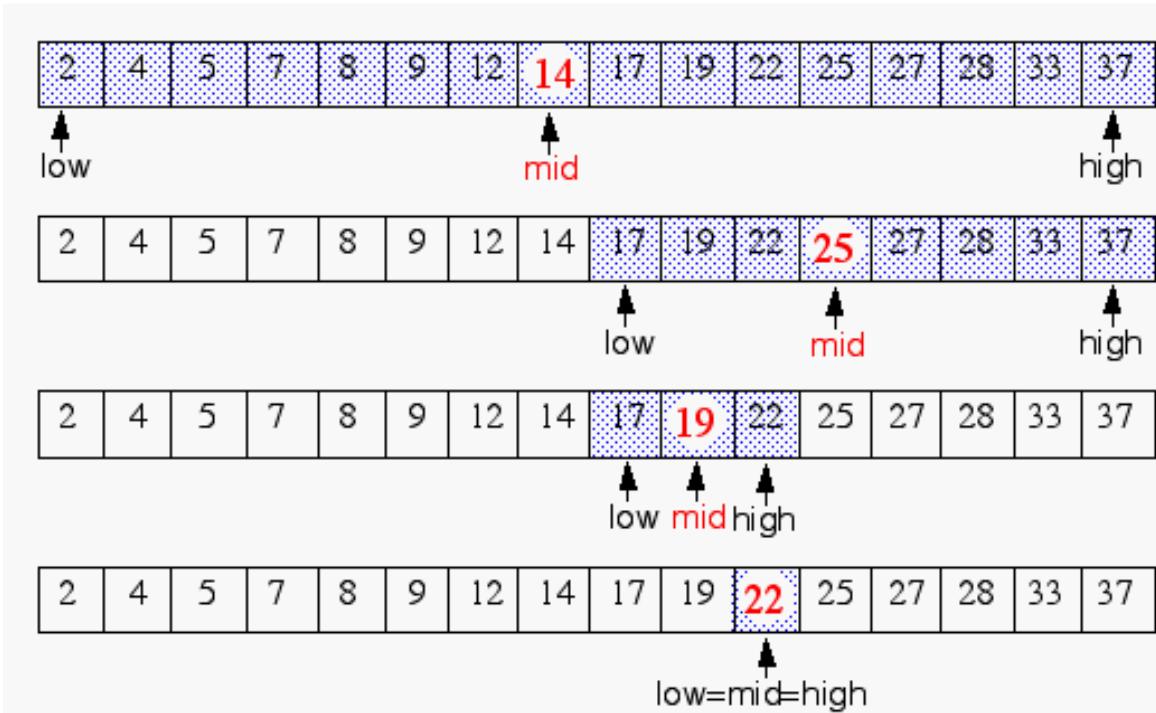
- ❑ Ordered array
- ❑ searching takes $O(\log n)$ time (binary search)
- ❑ inserting and removing takes $O(n)$ time
- ❑ application to look-up tables (frequent searches, rare insertions and removals)

- ❑ Apply binary search



Binary Searches

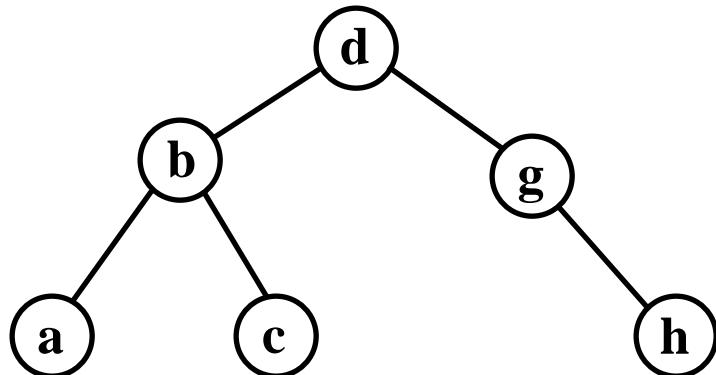
- ❑ narrow down the search range in stages
- ❑ “high-low” game
- ❑ `findElement(22)`





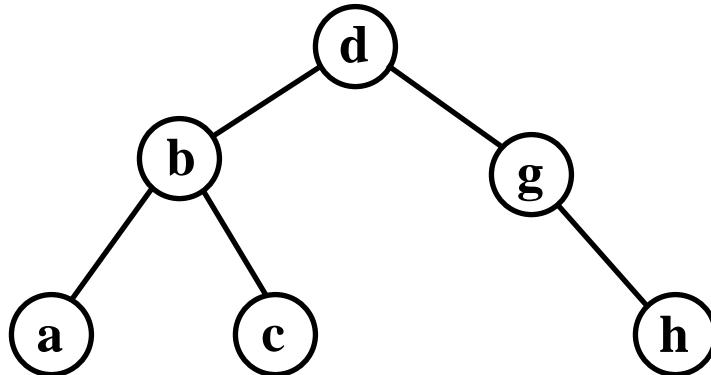
Recall Binary Search Trees...

- ❑ Implement a dictionary with a BST
 - A binary search tree is a binary tree T such that
 - each internal node stores **an item (k, e) of a dictionary**.
 - keys stored at nodes in the left subtree of v are less than or equal to k .
 - keys stored at nodes in the right subtree of v are greater than or equal to k .





Binary search trees

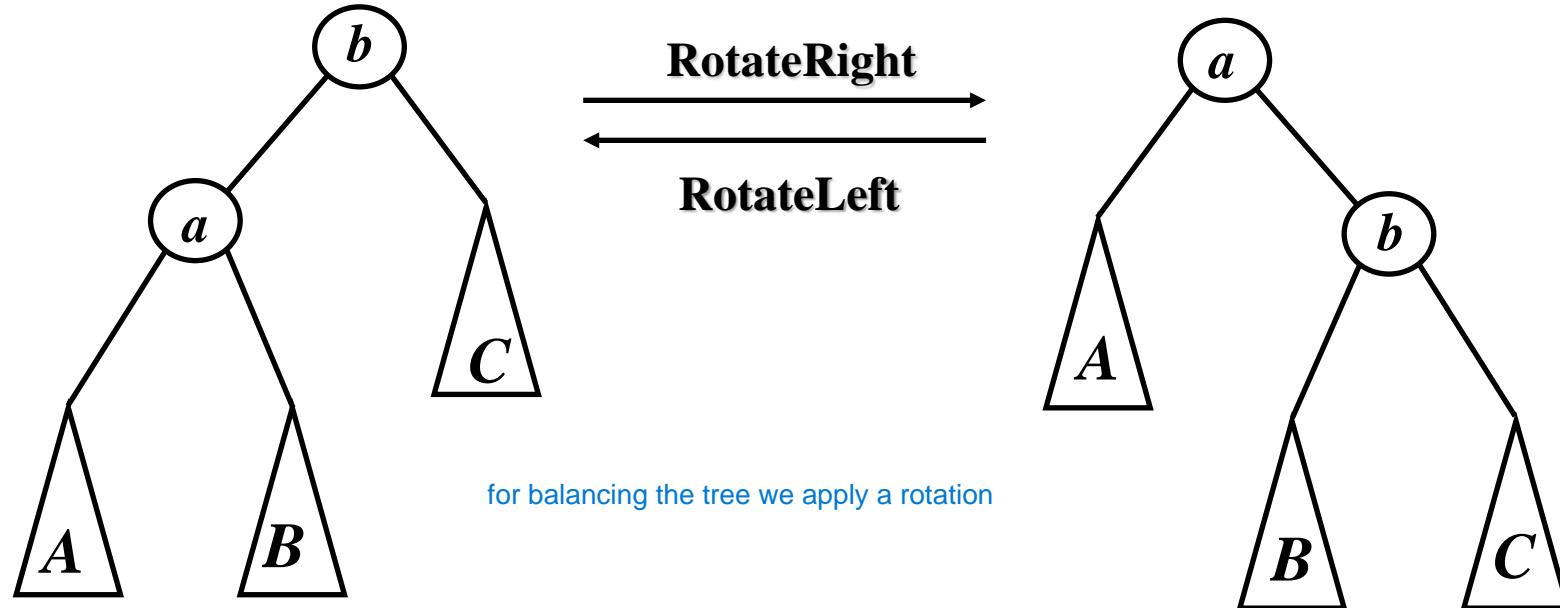


- BST Drawback: Sequence of insertions may lead to a linear list a, b, c, d, e, f
- AVL tree
 - Better performance, difficult to implement
- Splay tree
 - Good overall performance -- *amortized*



AVL tree - self-balancing binary search tree

- Several strategies have been devised to handle BST drawbacks, usually involving balancing operations to ensure that the tree has height $O(\log n)$
 - The most commonly used strategy is to perform *rotations* during the update operations





Splay tree

- A different strategy, called *splaying*, is used in "self-adjusting" search trees to guarantee an *amortized* time bound of $O(\log n)$;
 - the splay operation moves a specified node to the root via a sequence of rotations
 - Amortization is the partitioning of the total cost of a sequence of operations among the individual operations in that sequence; thus, an amortized time bound can be viewed as the average cost of the operations in a sequence



Splay tree

- The idea behind self-adjusting trees is to use a particular implementation of the splay operation to move to the root a node accessed by a FIND operation.
- If a node is accessed **often enough**, it will remain close to the root and will not contribute much to the total running time; an infrequently accessed node cannot contribute much to the total running time in any case
 - guarantee only **amortized** logarithmic time per operation
 - relatively simple to implement
 - do not require explicit **balance information** to be stored at nodes
 - can be shown to be **optimal** with respect to arbitrary access frequencies



Splay tree drawbacks

- they **restructure** the entire tree not only during updates but also while performing simple search operations
 - significant slowdown in practice in caching and paging environments
- during **any given operation** splay trees may perform a logarithmic number of rotations
- inefficient in implementing **higher dimensional search trees**
 - secondary data structures associated with each node of these higher dimensional trees, and the secondary data structure at any node depends on the set of keys stored in the sub-tree rooted at that node. Since the entire secondary data structure has to be recomputed during each rotation, the cost of performing a single rotation could increase from a constant to some super-linear function of the sub-tree size
- we do not have the **guarantee** that every operation will run quickly
 - we obtain bounds only on the total cost of the operations



Random Treaps

- treaps achieve essentially the same time bounds
- they do not require any explicit balance information
- the expected number of rotations performed is small for each operation
- simple to implement

- Skip lists similar benefits but alternative randomized data structure



Approach for randomized search trees

- If n elements are inserted in random order into a binary search tree, the expected depth is $1.39 \log n$
- Goal:
 - the search tree has the structure that would result if elements were inserted in the order of their priorities (property 1)
- Idea:
 - Each element x is assigned a priority chosen uniformly at random
$$\text{prio}(x) \in R$$



Treaps: tree + heap

- **Definition:** A treap is a binary tree
- Each node contains one element x with $\text{key}(x) \in U$ and $\text{prio}(x) \in R$
- The following properties hold:
 - Search tree property
 - For each element x :
 - elements y in the left subtree of x satisfy: $\text{key}(y) < \text{key}(x)$
 - elements y in the right subtree of x satisfy : $\text{key}(y) > \text{key}(x)$
 - Heap property
 - For all elements x,y :
 - If y is a child of x , then $\text{prio}(y) > \text{prio}(x)$.
 - All priorities are pairwise distinct.



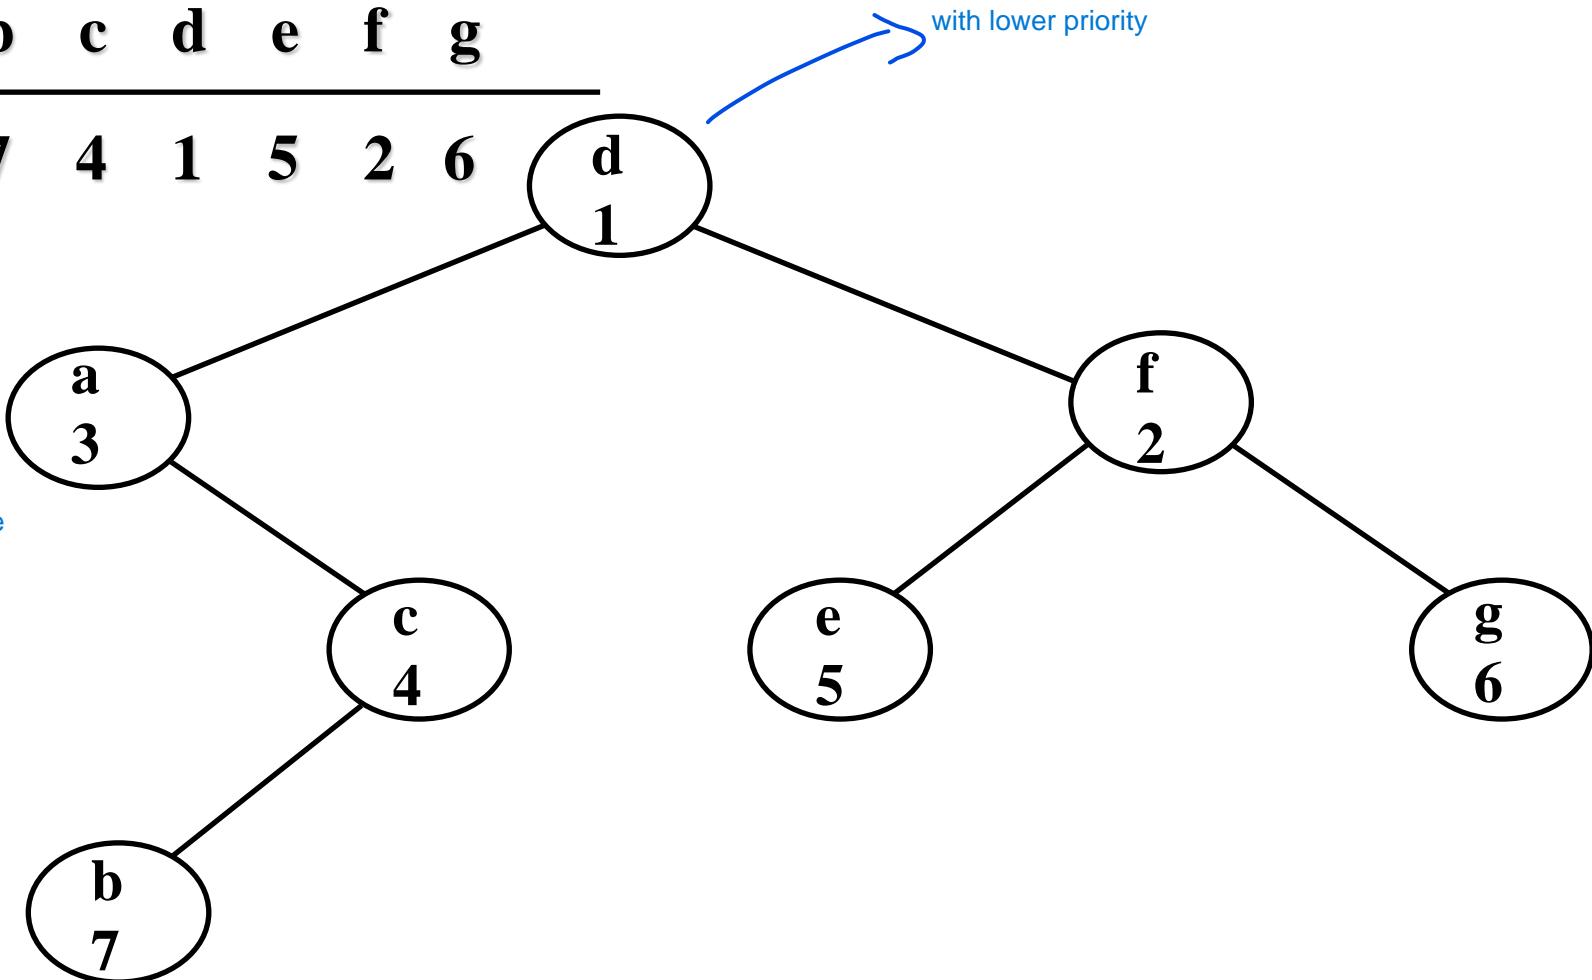
Example

key	a	b	c	d	e	f	g
Priority	3	7	4	1	5	2	6

Priority
3

nota che nel sottoalbero sx della root tutte le chiavi sono < di d.

e vale anche l'altra prop.
che i nodi figli hanno priorità maggiore di quella del genitore





Treap uniqueness

- **Lemma:** For elements x_1, \dots, x_n with $\text{key}(x_i)$ and $\text{prio}(x_i)$, there exists a unique treap.
- **Proof by induction on n:**
- $n=1$: ok
- $n>1$: *in that case we're focusing on the root and the children*
 - The root has the smallest priority (k_1, p_1)
 - All the elements y with $\text{key}(y) < k_1$ go on the left subtree $\leq n-1$
 - All the elements y with $\text{key}(y) > k_1$ go on the right subtree $\leq n-1$



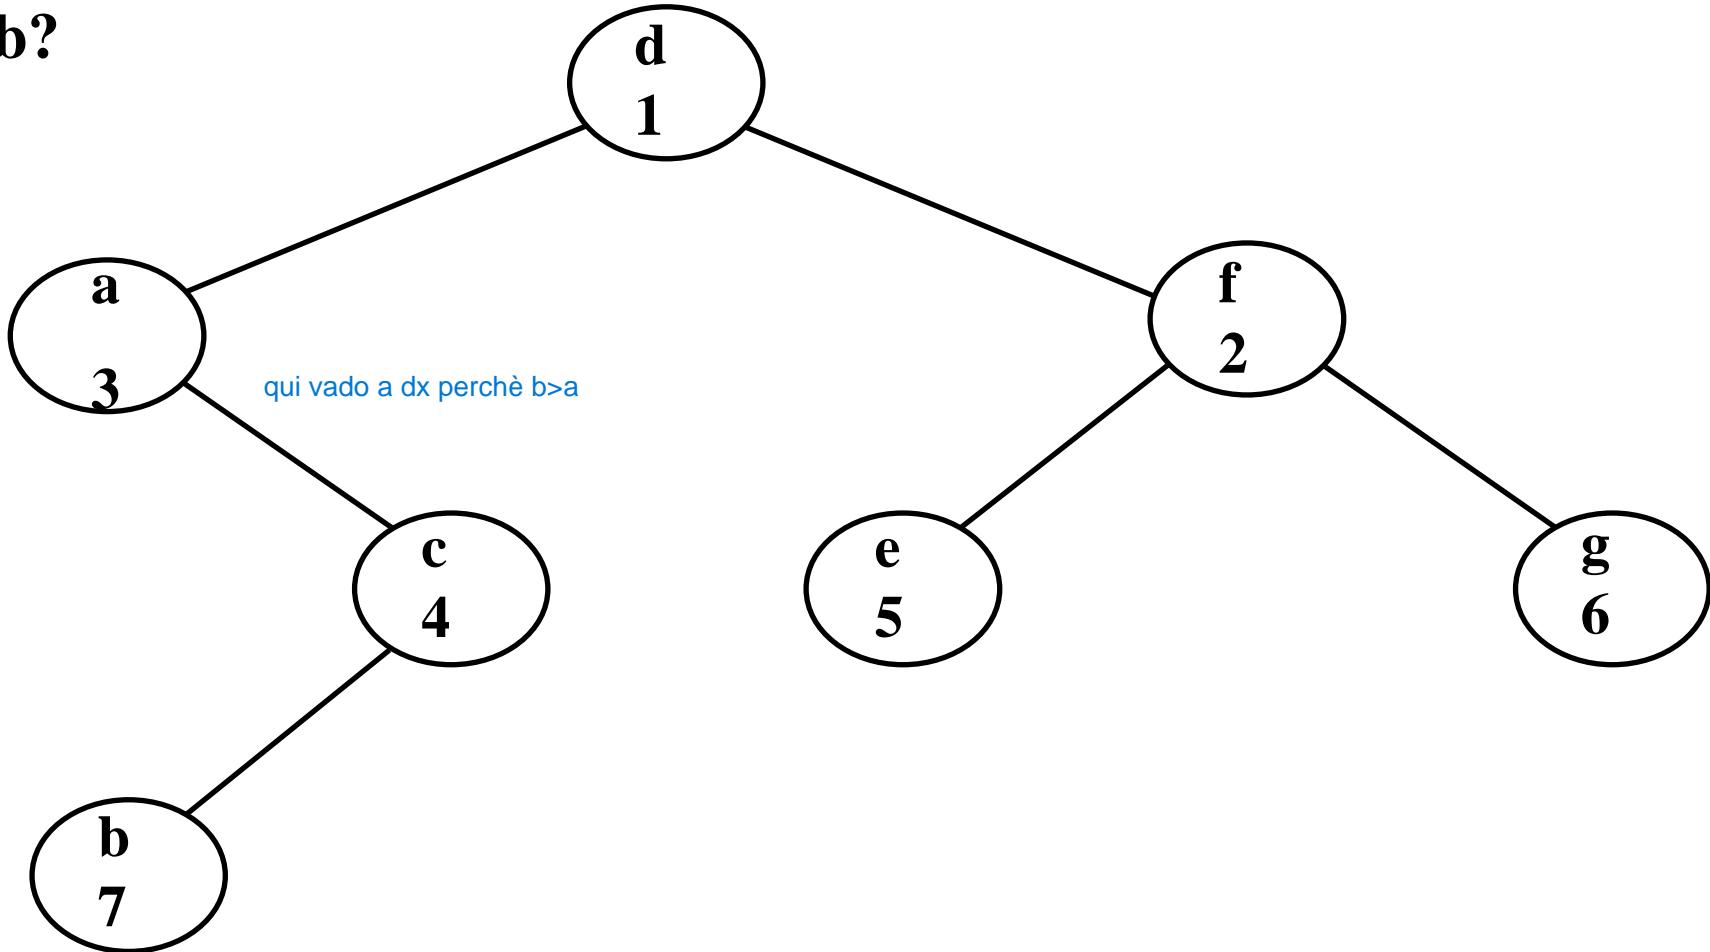
Treap structure

- the search tree has the structure that would result if elements were inserted in the order of their priorities
- **Proof by induction on n:**
- $n=1$: ok
- $n>1$:
 - The element with the smallest priority is in the root whatever it is the insertion order (k_1, p_1)
 - All y such that $\text{key}(y) < k_1$ go in the left subtree
 - All y such that $\text{key}(y) > k_1$ go in the right subtree



Search for an element: key

b?





Search for element with key k

```
1  $v := \text{root};$ 
2 while  $v \neq \text{nil}$  do
3   case  $\text{key}(v) = k$  : stop; "element found"
  (successful search)
4    $\text{key}(v) < k$  :  $v := \text{RightChild}(v);$ 
5    $\text{key}(v) > k$  :  $v := \text{LeftChild}(v);$ 
6 endcase;
7 endwhile;
8 "element not found" (unsuccessful search)
```

Running time: $O(\# \text{ elements on the search path})$



Analysis of the search path

Elements x_1, \dots, x_n

- x_i has i -th smallest key
- $\text{key}(x_1) < \text{key}(x_2) < \text{key}(x_3) < \dots < \text{key}(x_n)$

Let M be a subset of the elements.

$P_{\min}(M)$ = element in M with lowest priority

Let's find out x_m

Lemma:

- a) Let $i < m$. x_i is ancestor of x_m iff $P_{\min}(\{x_i, \dots, x_m\}) = x_i$
- b) Let $m < i$. x_i is ancestor of x_m iff $P_{\min}(\{x_m, \dots, x_i\}) = x_i$



Analysis of the search path

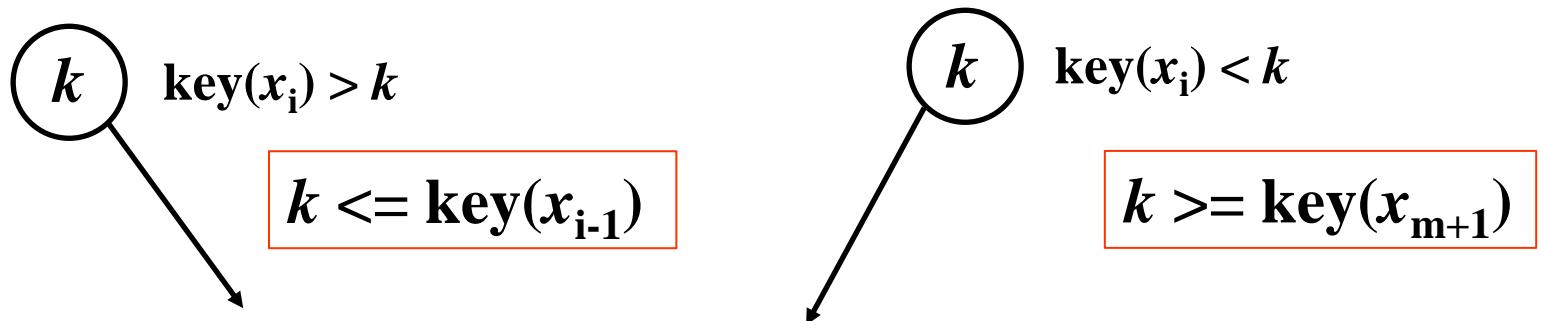
Proof: a) Use of property 1. Elements are inserted in order of increasing priorities.

“ \Leftarrow ” $P_{\min}(\{x_1, \dots, x_m\}) = x_i \Rightarrow x_i$ is inserted first among $\{x_1, \dots, x_m\}$.

When x_i is inserted, the tree contains only keys k with

$k < \text{key}(x_i)$ or $k > \text{key}(x_m)$

$\text{key}(x_i) < \dots < \text{key}(x_m)$





Analysis of the search path

Proof: a) (Let $i < m$. x_i is ancestor of x_m iff $P_{\min}(\{x_i, \dots, x_m\}) = x_i$)

“ \Rightarrow ” Let $x_j = P_{\min}(\{x_i, \dots, x_m\})$. Show: $x_i = x_j$

Suppose: $x_i \neq x_j$

Use property 1, considering the search path when x_j is inserted

As before any x_l traverse the same path as x_j

imply that x_j is ancestor of x_l

Case 1: $x_j = x_m$

Case 2: $x_j \neq x_m$

Part b) follows analogously.



Analysis of the ‘Search’ operation

Let T be a treap with elements x_1, \dots, x_n x_i has i -th smallest key

n -th Harmonic number:

$$H_n = \sum_{k=1}^n 1/k$$

$$H_n = \ln n + O(1)$$

Lemma:

1. Successful search: The expected number of nodes on the path to x_m is $H_m + H_{n-m+1} - 1$.
2. Unsuccessful search : Let m be the number of keys that are smaller than the search key k . The expected number of nodes on the search path is $H_m + H_{n-m}$.



Analysis of the ‘Search’ operation

Proof: Part 1

$$X_{m,i} = \begin{cases} 1 & x_i \text{ is ancestor of } x_m \\ 0 & \text{otherwise} \end{cases}$$

$X_m = \# \text{ nodes on the path from the root to } x_m (\text{incl. } x_m)$

$$X_m = 1 + \sum_{i < m} X_{m,i} + \sum_{i > m} X_{m,i}$$

in case i is before m in case i is after m

$$E[X_m] = 1 + E\left[\sum_{i < m} X_{m,i}\right] + E\left[\sum_{i > m} X_{m,i}\right]$$



Analysis of the ‘Search’ operation

$i < m :$

$$E[X_{m,i}] = \text{Prob}[x_i \text{ is ancestor of } x_m] = 1/(m-i+1)$$

All elements in $\{x_i, \dots, x_m\}$ have the same probability of being the one with the smallest priority

$$\text{Prob}[P_{\min}(\{x_i, \dots, x_m\}) = x_i] = 1/(m-i+1)$$

$i > m :$

$$E[X_{m,i}] = 1/(i-m+1)$$



Analysis of the ‘Search’ operation

$$\begin{aligned} E[X_m] &= 1 + \sum_{i < m} \frac{1}{m-i+1} + \sum_{i > m} \frac{1}{i-m+1} \\ &= 1 + \frac{1}{m} + \dots + \frac{1}{2} + \frac{1}{2} + \dots + \frac{1}{n-m+1} \\ &= H_m + H_{n-m+1} - 1 \end{aligned}$$

so the searching is going to work with a complexity of $\log n$ (due to the fact we're considering binary tree)

Part 2 follows analogously



Inserting a new element x

1. Choose $\text{prio}(x)$.
2. Search for the position of x in the tree.

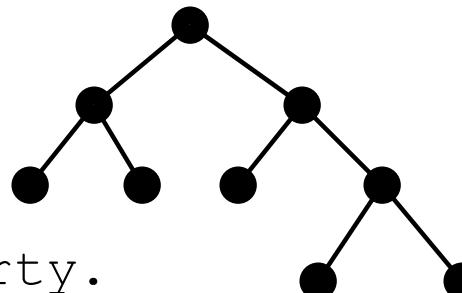
constant time

3. Insert x as a leaf.
4. Restore the heap property.

```
while  $\text{prio}(\text{parent}(x)) > \text{prio}(x)$  do
    if  $x$  is left child then  $\text{RotateRight}(\text{parent}(x))$ 
    else  $\text{RotateLeft}(\text{parent}(x));$ 
    endif
endwhile;
```

costant time

logn cost

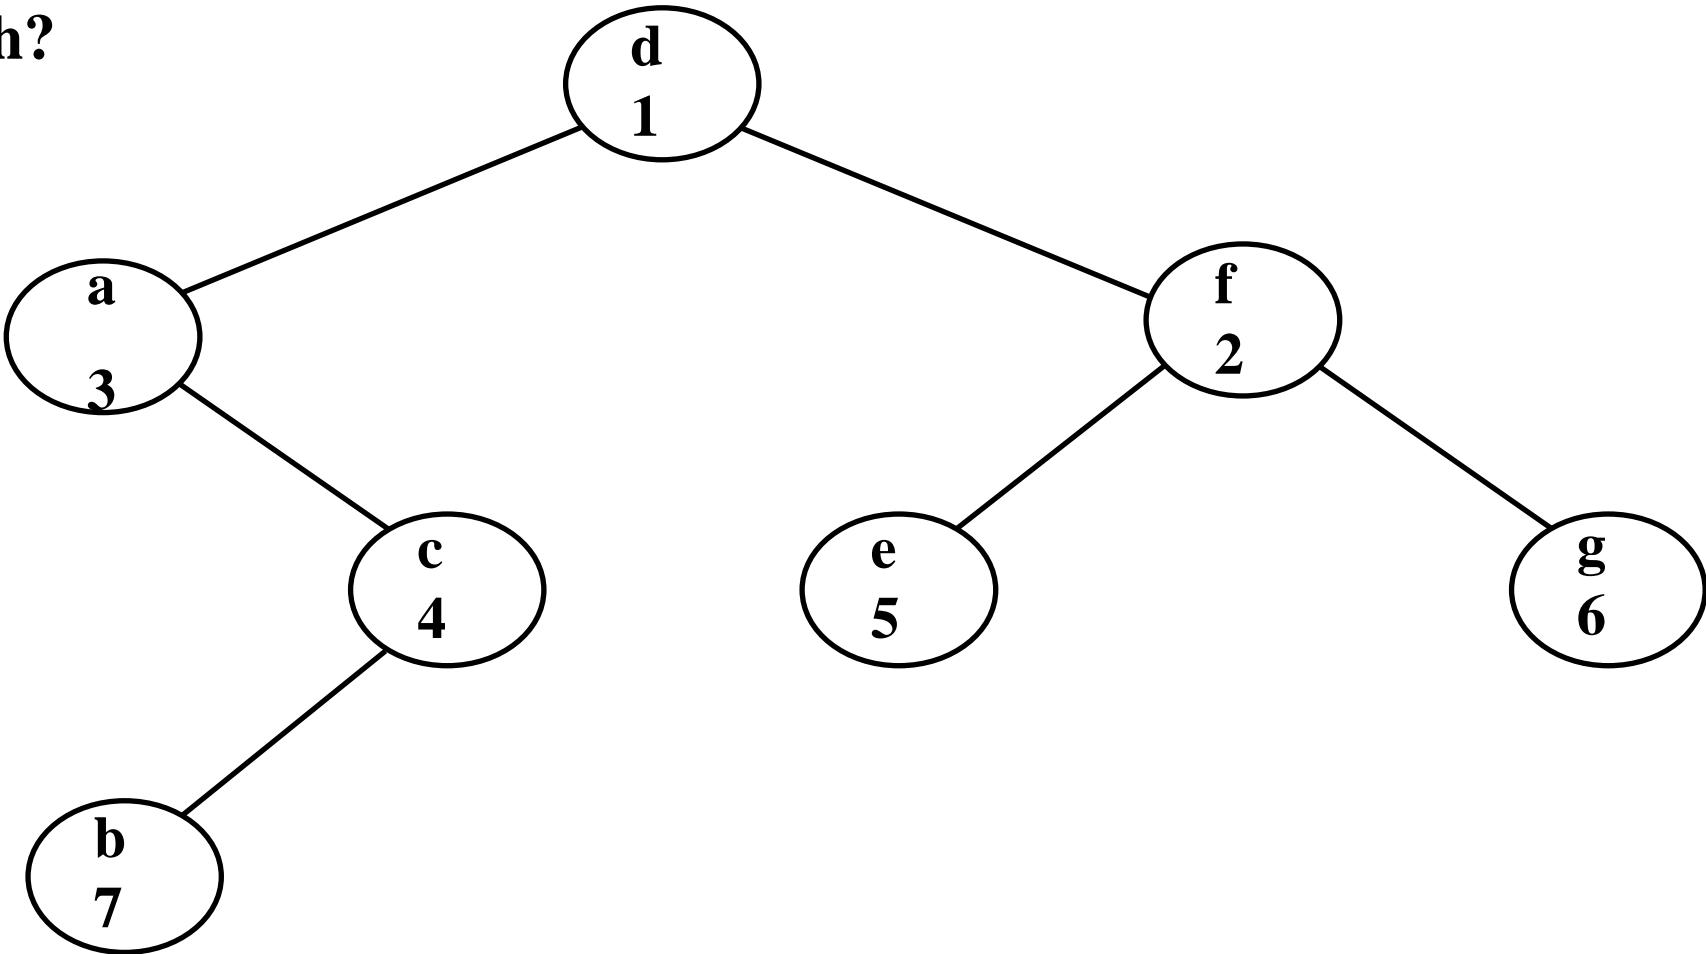


the rotation costs in the order of logn.



Insert an element

h?

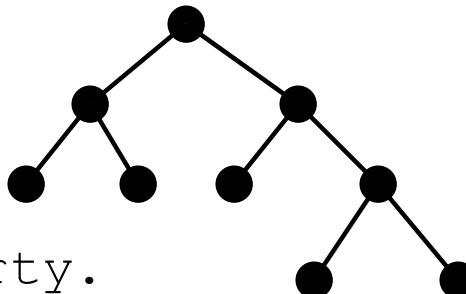




Inserting a new element x

1. Choose $\text{prio}(x)$.
2. Search for the position of x in the tree.

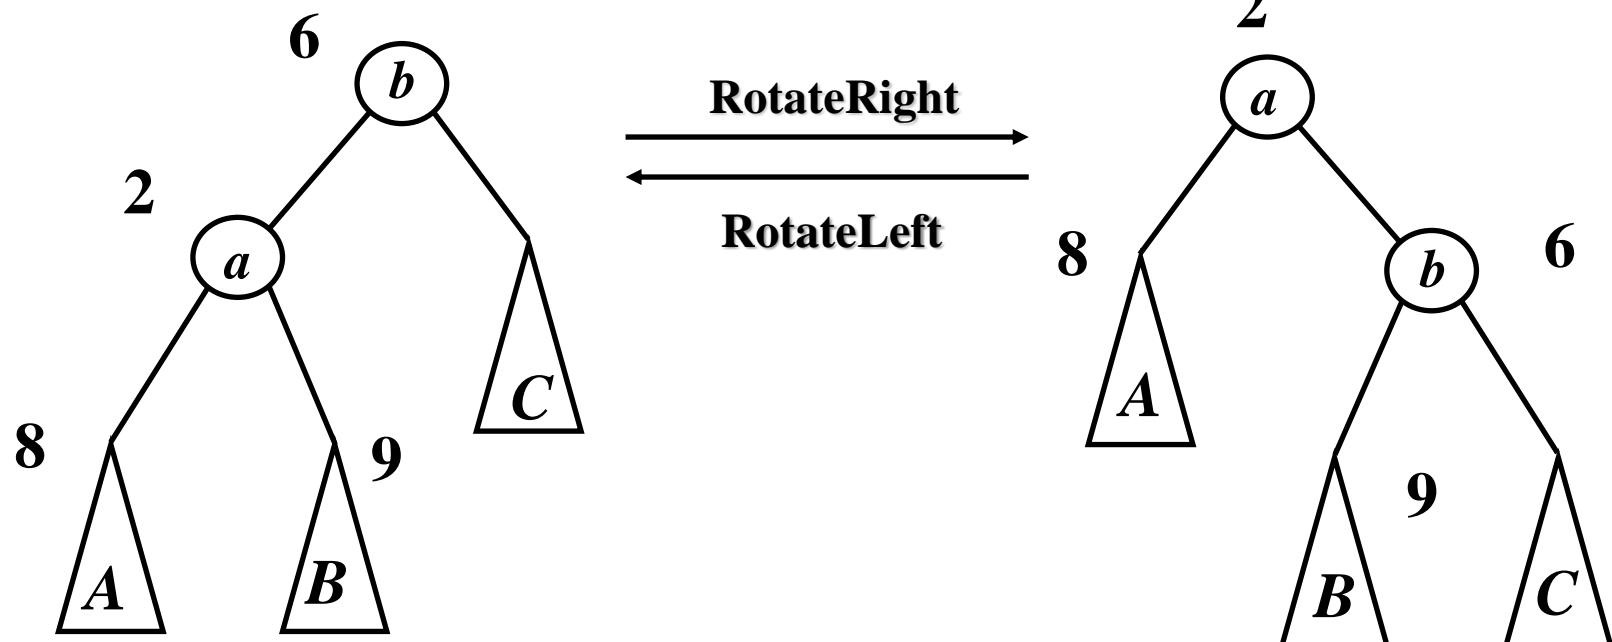
3. Insert x as a leaf.
4. Restore the heap property.



```
while  $\text{prio}(\text{parent}(x)) > \text{prio}(x)$  do
    if  $x$  is left child then  $\text{RotateRight}(\text{parent}(x))$ 
    else  $\text{RotateLeft}(\text{parent}(x))$ ;
    endif
endwhile;
```



Rotations



- The rotations maintain the search tree property and restore the heap property.



Deleting an element x

1. Find x in the tree.

2. **while** x is not a leaf **do**

u := child with smaller priority;

if *u* is left child **then** RotateRight(*x*)

else RotateLeft(*x*);

endif;

endwhile;

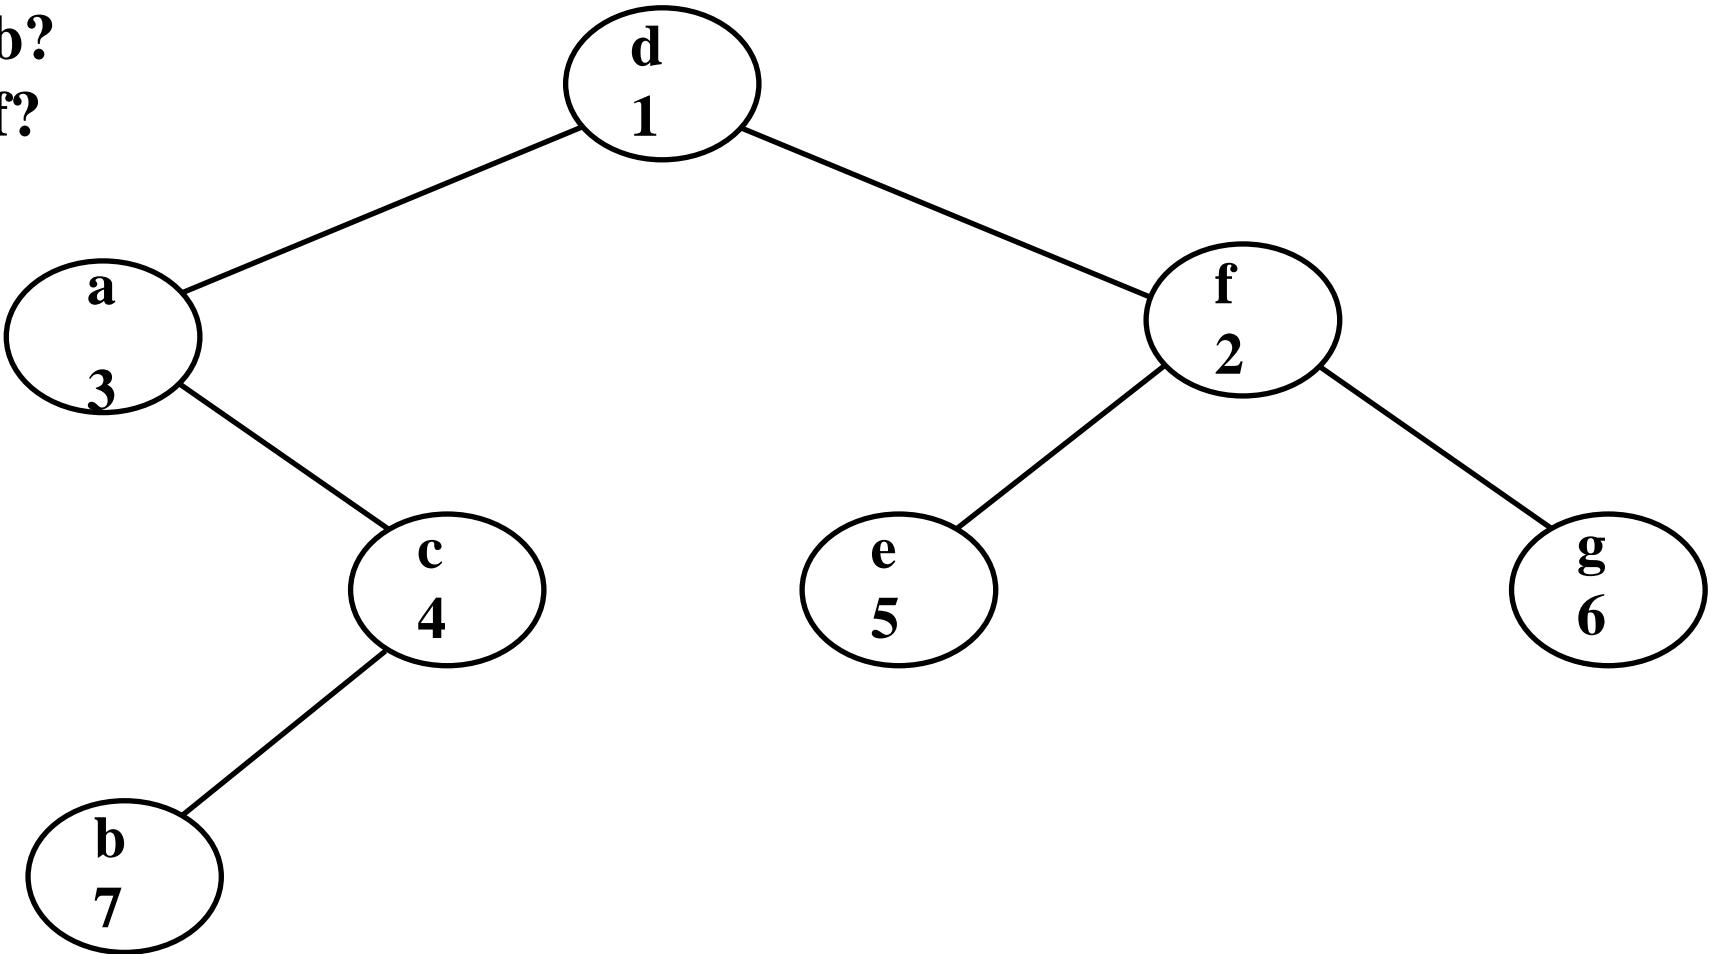
3. Delete x;

if the element is a leaf, we can remove directly the node in a constant time. otherwise we have to do a specific number of rotations before deleting it



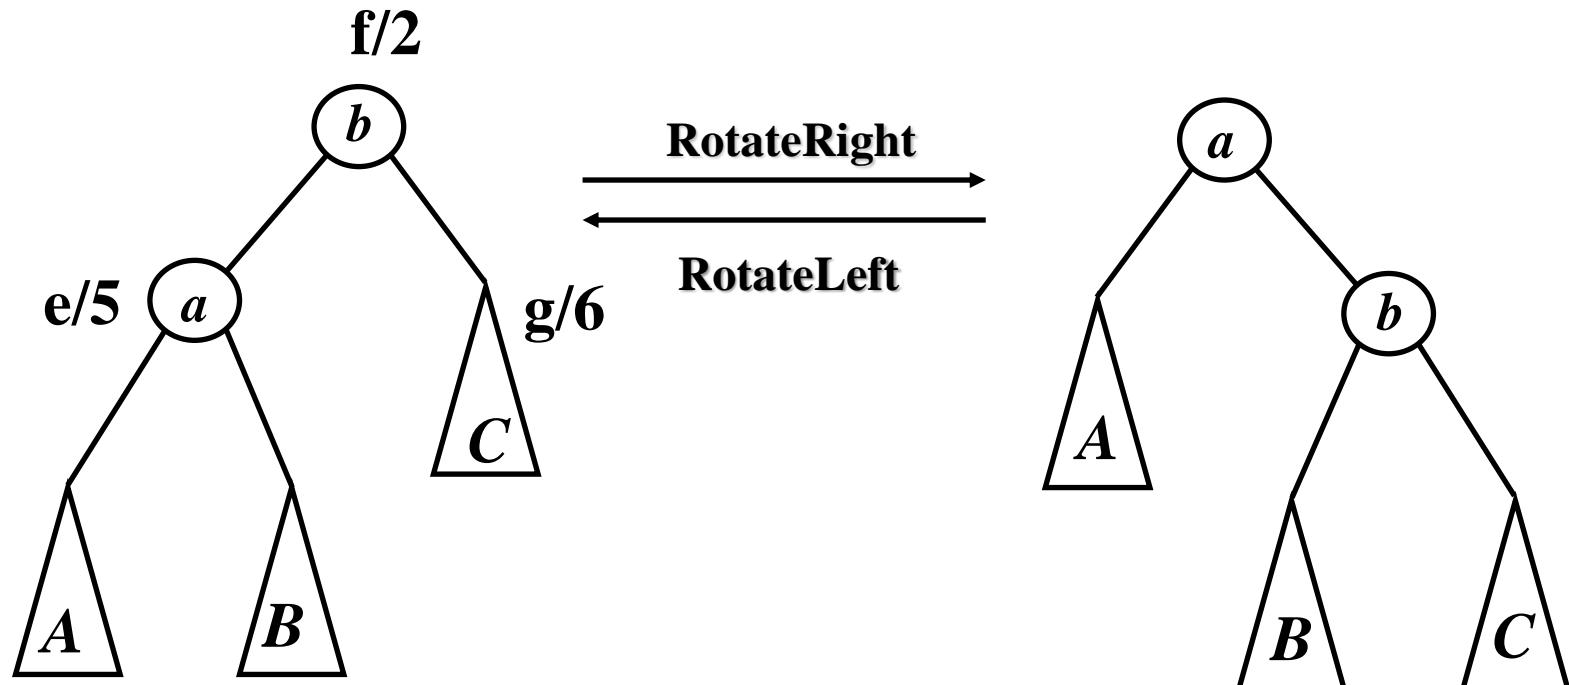
Delete an element

b?
f?





Rotations





Analysis of ‘Insert’ and ‘Delete’ operations

Lemma: The expected running time of insert and delete operations is $O(\log n)$. The expected number of rotations is 2.

Proof: Analysis of insert (delete is the inverse operation)

rotations = depth of x after being inserted as a leaf (1)

- depth of x after the rotations (2)

Let $x = x_m$.

(2) Expected depth is $H_m + H_{n-m+1} - 1$.

(1) Expected depth is $H_{m-1} + H_{n-m} + 1$.

The tree contains $n-1$ elements, $m-1$ of them being smaller.

rotations = $H_{m-1} + H_{n-m} + 1 - (H_m + H_{n-m+1} - 1) < 2$



Extended set of operations

n = number of elements in treap T .

- $\text{Minimum}(T)$: Return the smallest key. $O(\log n)$
- $\text{Maximum}(T)$: Return the largest key. $O(\log n)$
- $\text{List}(T)$: Output elements of S in increasing order. $O(n)$
- $\text{Union}(T_1, T_2)$: Merge T_1 and T_2 .
Condition: $\forall x_1 \in T_1, x_2 \in T_2: \text{key}(x_1) < \text{key}(x_2)$
- $\text{Split}(T, k, T_1, T_2)$: Split T into T_1 and T_2 .
 $\forall x_1 \in T_1, x_2 \in T_2: \text{key}(x_1) \leq k \text{ and } k < \text{key}(x_2)$



The ‘Split’ operation

$\text{Split}(T, k, T_1, T_2)$: Split T into T_1 and T_2 .

$$\forall x_1 \in T_1, x_2 \in T_2: \text{key}(x_1) \leq k \text{ and } k < \text{key}(x_2)$$

W.l.o.g. key k is not in T .

Otherwise delete the element with key k and re-insert it into T_1 after the split operation.

1. Generate a new element x with $\text{key}(x)=k$ and $\text{prio}(x)=-\infty$. (so the root \rightarrow minimum priority)
2. Insert x into T .
3. Delete the new root. The left subtree is T_1 , the right subtree is T_2 .



The ‘Union’ operation

$\text{Union}(T_1, T_2)$: Merge T_1 and T_2 .

Condition: $\forall x_1 \in T_1, x_2 \in T_2: \text{key}(x_1) < \text{key}(x_2)$

1. Determine key k with $\text{key}(x_1) < k < \text{key}(x_2)$ for all $x_1 \in T_1$ and $x_2 \in T_2$. 
it costs $\log n$
2. Generate element x with $\text{key}(x)=k$ and $\text{prio}(x) = -\infty$.
3. Generate treap T with root x , left subtree T_1 and right subtree T_2 .
4. Delete x from T .



Analysis

- **Lemma:** The expected running time of the operations Union and Split is $O(\log n)$.

Now there is a problem. -> when we're generating with a random algorithm priority we have to consider that the child of the parent must have a priority < than the parent's priority. -> we have to manage the random algorithm taking care also of this problem



Implementation

- Priorities from [0,1)
- Priorities are used only when two elements are compared to find out which of them has the higher priority. stiamo considerando le priorità rappresentabili con un certo numero di bits
- In case of equality, extend both priorities by bits chosen uniformly at random until two corresponding bits differ.

$$p_1 = 0,010111001$$

$$p_2 = 0,010111001$$

$$p_1 = 0,010111001\textcolor{red}{0}11$$

$$p_2 = 0,010111001\textcolor{red}{0}10$$

genero i primi bit rossi, se sono uguali ne ripeto altri (sempre randomicamente) fino ad ottenere due priorità diverse



Skip lists

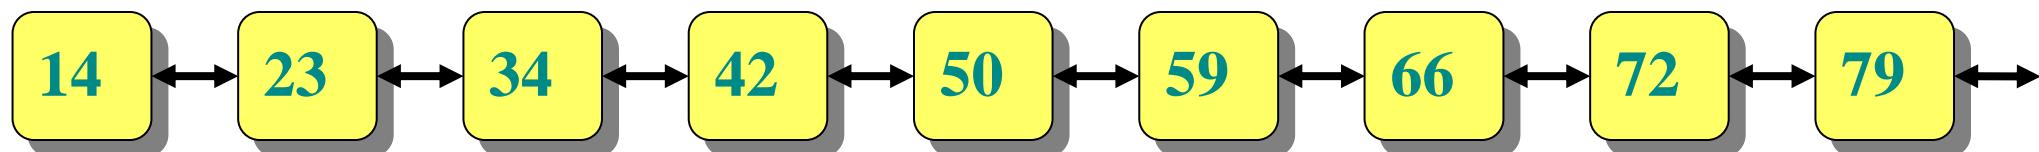
- Simple randomized dynamic search structure
 - Invented by William Pugh in 1989
 - Easy to implement
- Maintains a dynamic set of n elements in $O(\lg n)$ time per operation in expectation and *with high probability*
 - Strong guarantee on tail of distribution of $T(n)$
 - $O(\lg n)$ “almost always”



One linked list

Start from simplest data structure:
(sorted) linked list

- Searches take $\Theta(n)$ time in worst case
- How can we speed up searches?

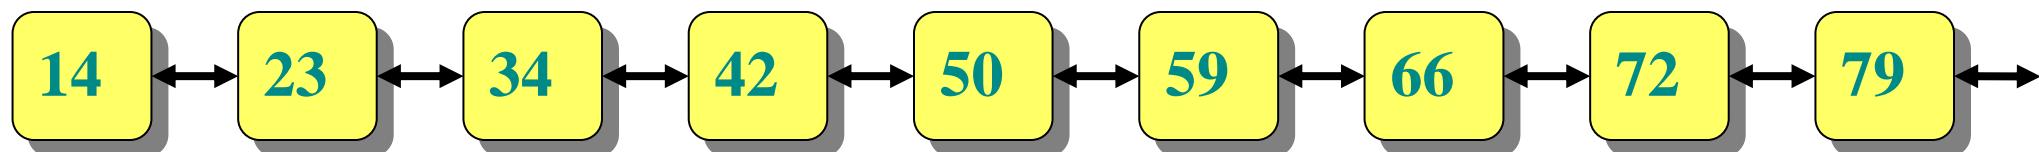




Two linked lists

Suppose we had *two* sorted linked lists
(on subsets of the elements)

- Each element can appear in one or both lists
- How can we speed up searches?



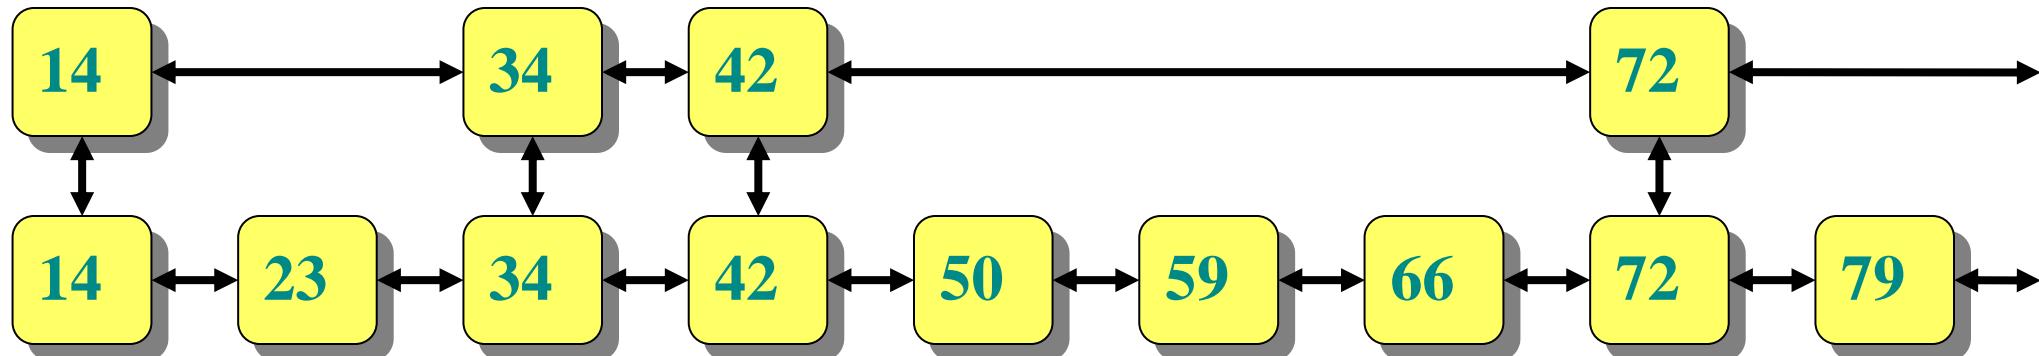


Two linked lists as a subway

IDEA: Express and local subway lines
(à la New York City 7th Avenue Line)

- Express line connects a few of the stations
- Local line connects all stations
- Links between lines at common stations

sta prendendo a caso $\log n$ elementi dalla lista di sotto e li sta mettendo nella lista di sopra che in qualche modo è collegata alla seconda. Così se sono fortunato posso trovare l'elemento che sto cercando subito pagando solo $\log n$ (che è minore di n), se invece non la trovo lì, posso cercarla sotto visto che ho dei collegamenti ma comunque spendo meno di n

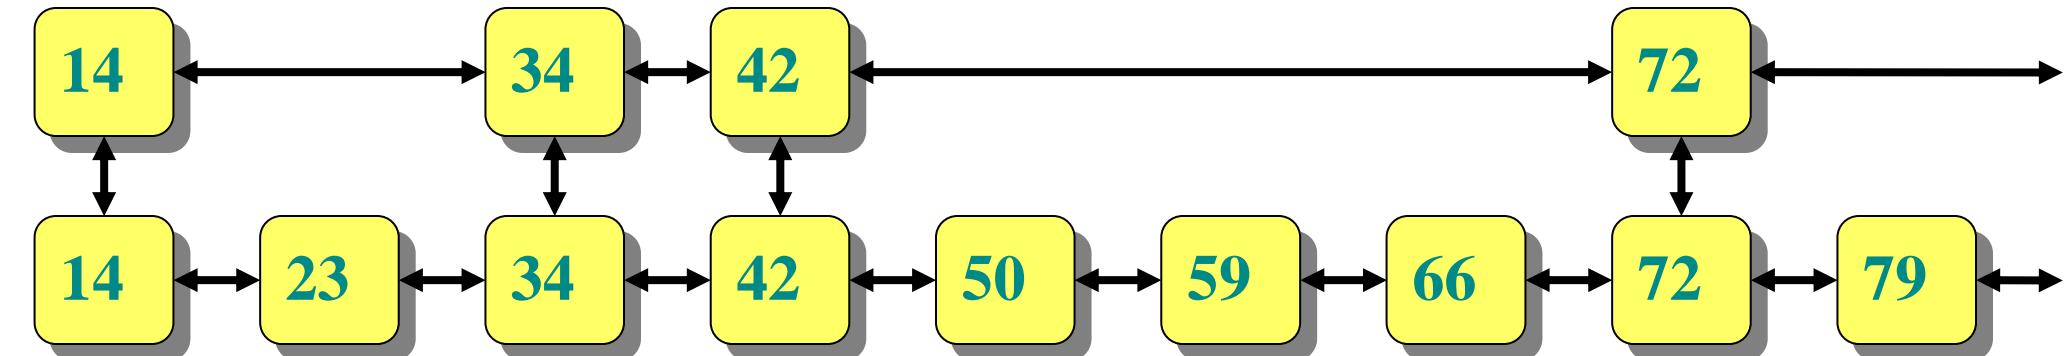




Searching in two linked lists

SEARCH(x):

- Walk right in top linked list (L_1) until going right would go too far
- Walk down to bottom linked list (L_2)
- Walk right in L_2 until element found (or not)

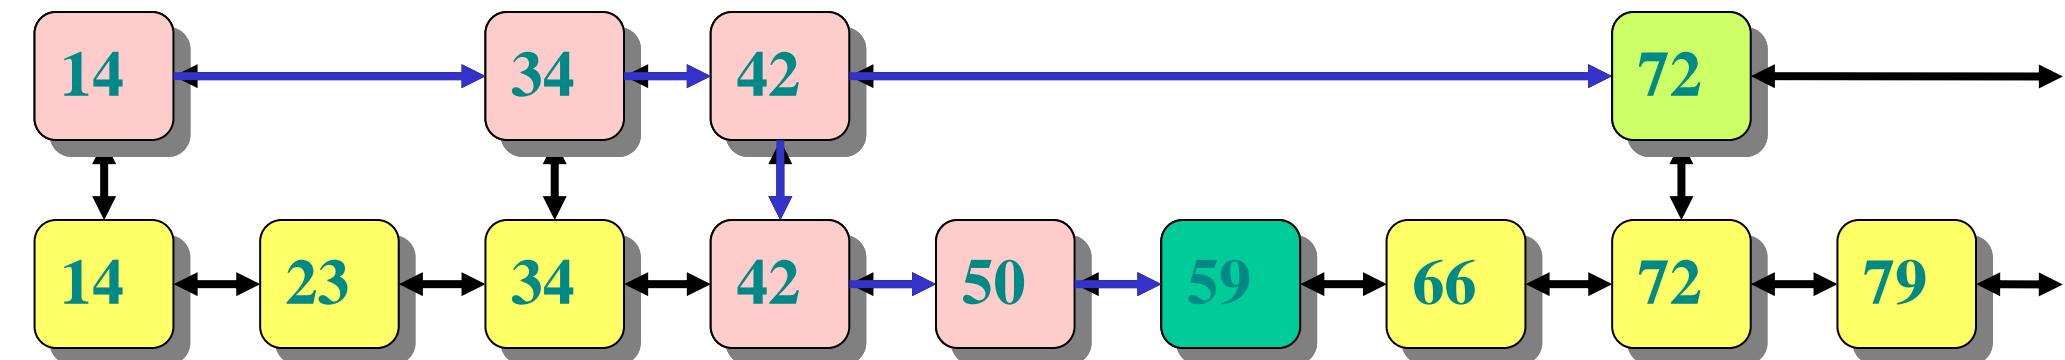




Searching in two linked lists

EXAMPLE: SEARCH(59)

Too far:
 $59 < 72$

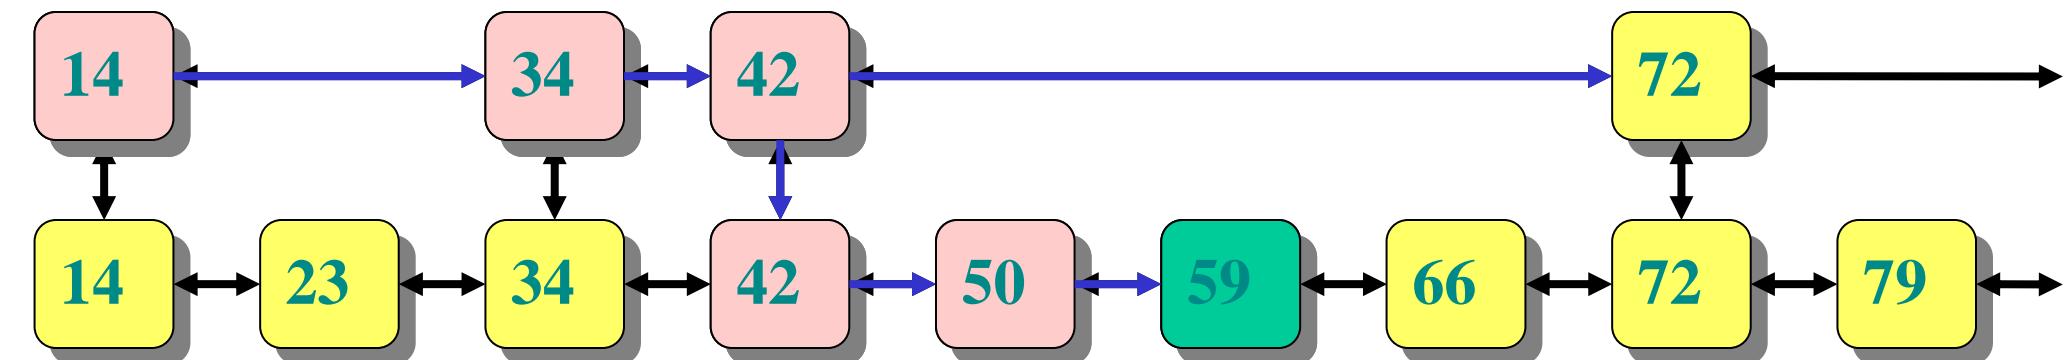




Design of two linked lists

QUESTION: Which nodes should be in L_1 ?

- In a subway, the “popular stations”
- Here we care about *worst-case performance*
- **Best approach:** Evenly space the nodes in L_1
- But *how many nodes* should be in L_1 ?



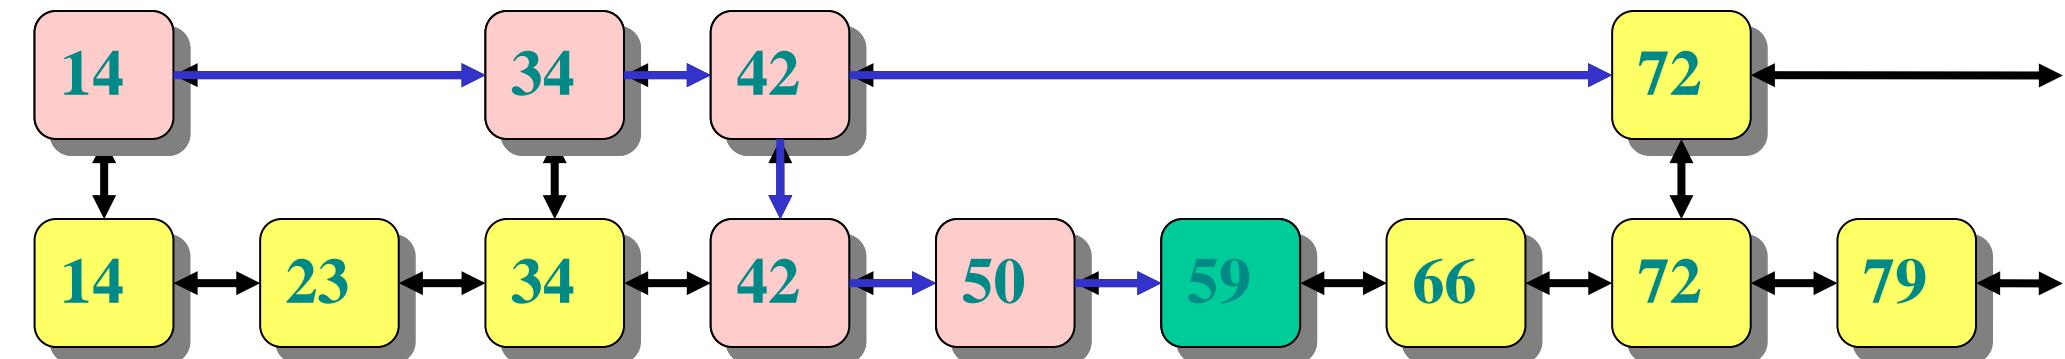


Analysis of two linked lists

ANALYSIS:

- Search cost is roughly $|L_1| + \frac{|L_2|}{|L_1|}$
- Minimized (up to constant factors) when terms are equal

$$|L_1|^2 = |L_2| = n \Rightarrow |L_1| = \sqrt{n}$$





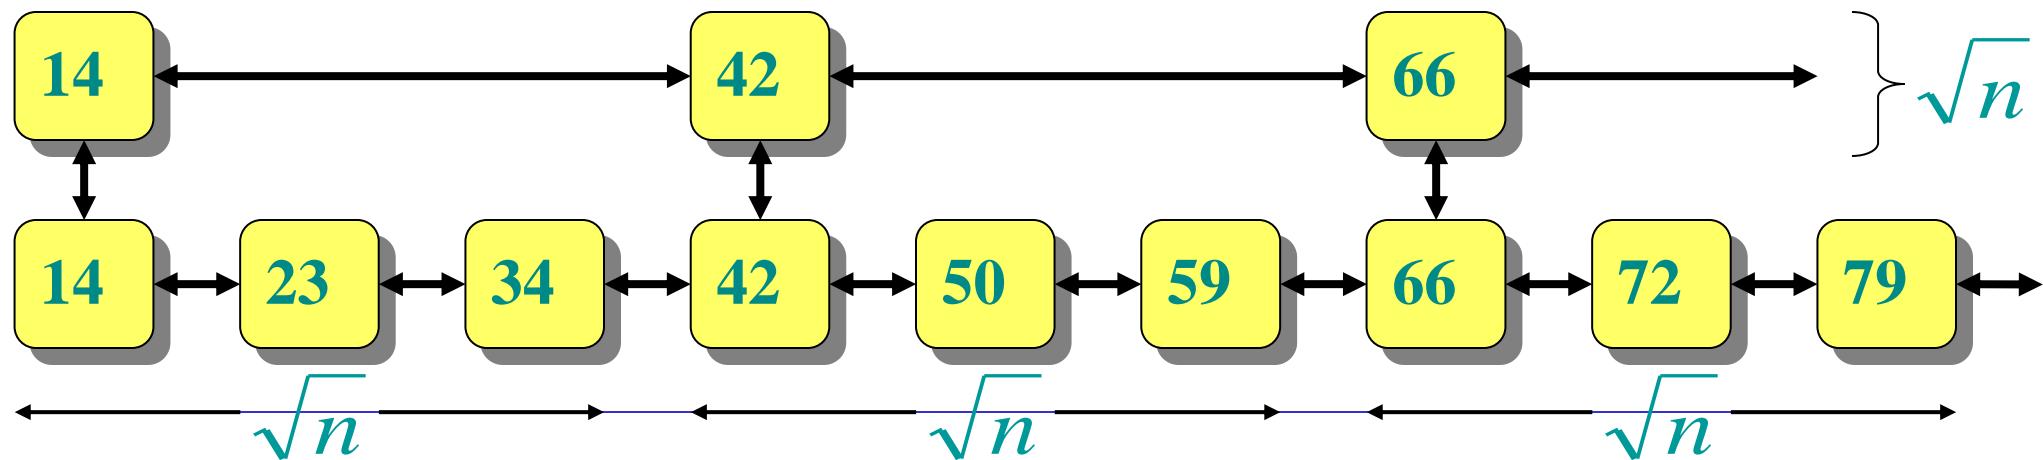
Analysis of two linked lists

ANALYSIS:

- $|L_1| = \sqrt{n}$ $|L_2| = n$

- Search cost is roughly

$$|L_1| + \frac{|L_2|}{|L_1|} = \sqrt{n} + \frac{n}{\sqrt{n}} = 2\sqrt{n}$$

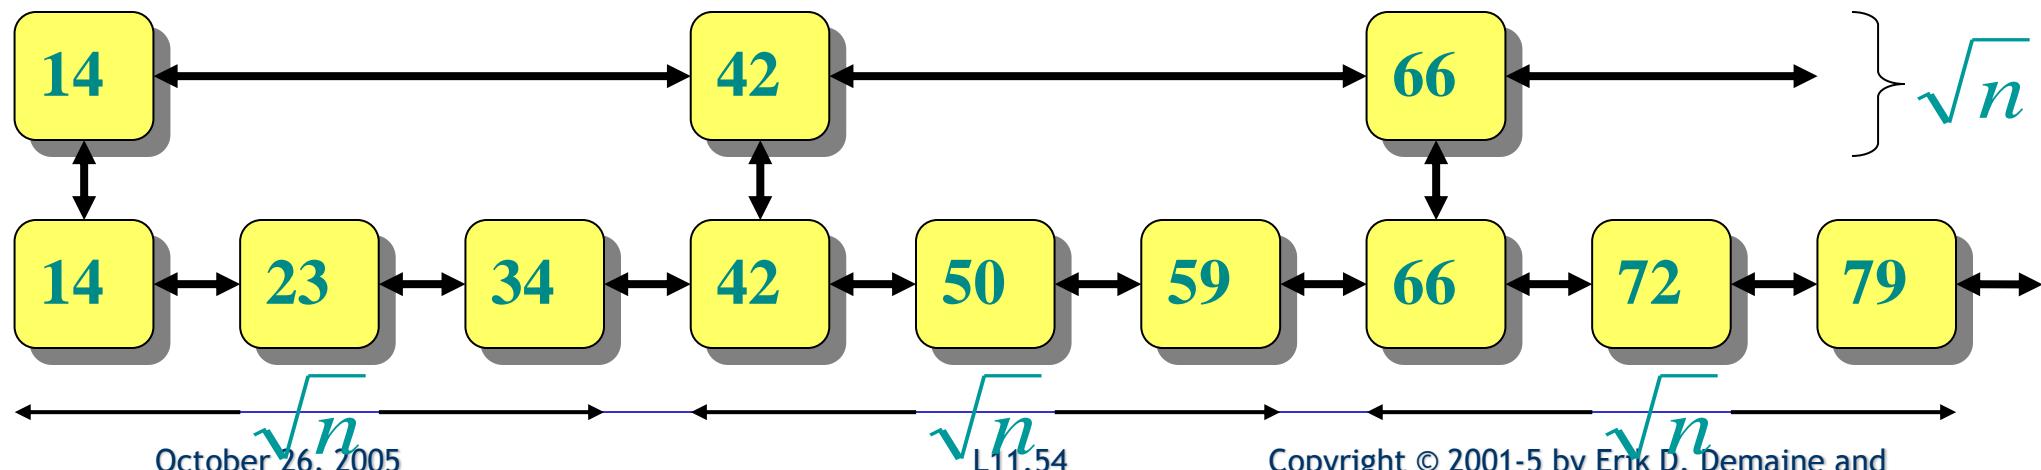




More linked lists

What if we had more sorted linked lists?

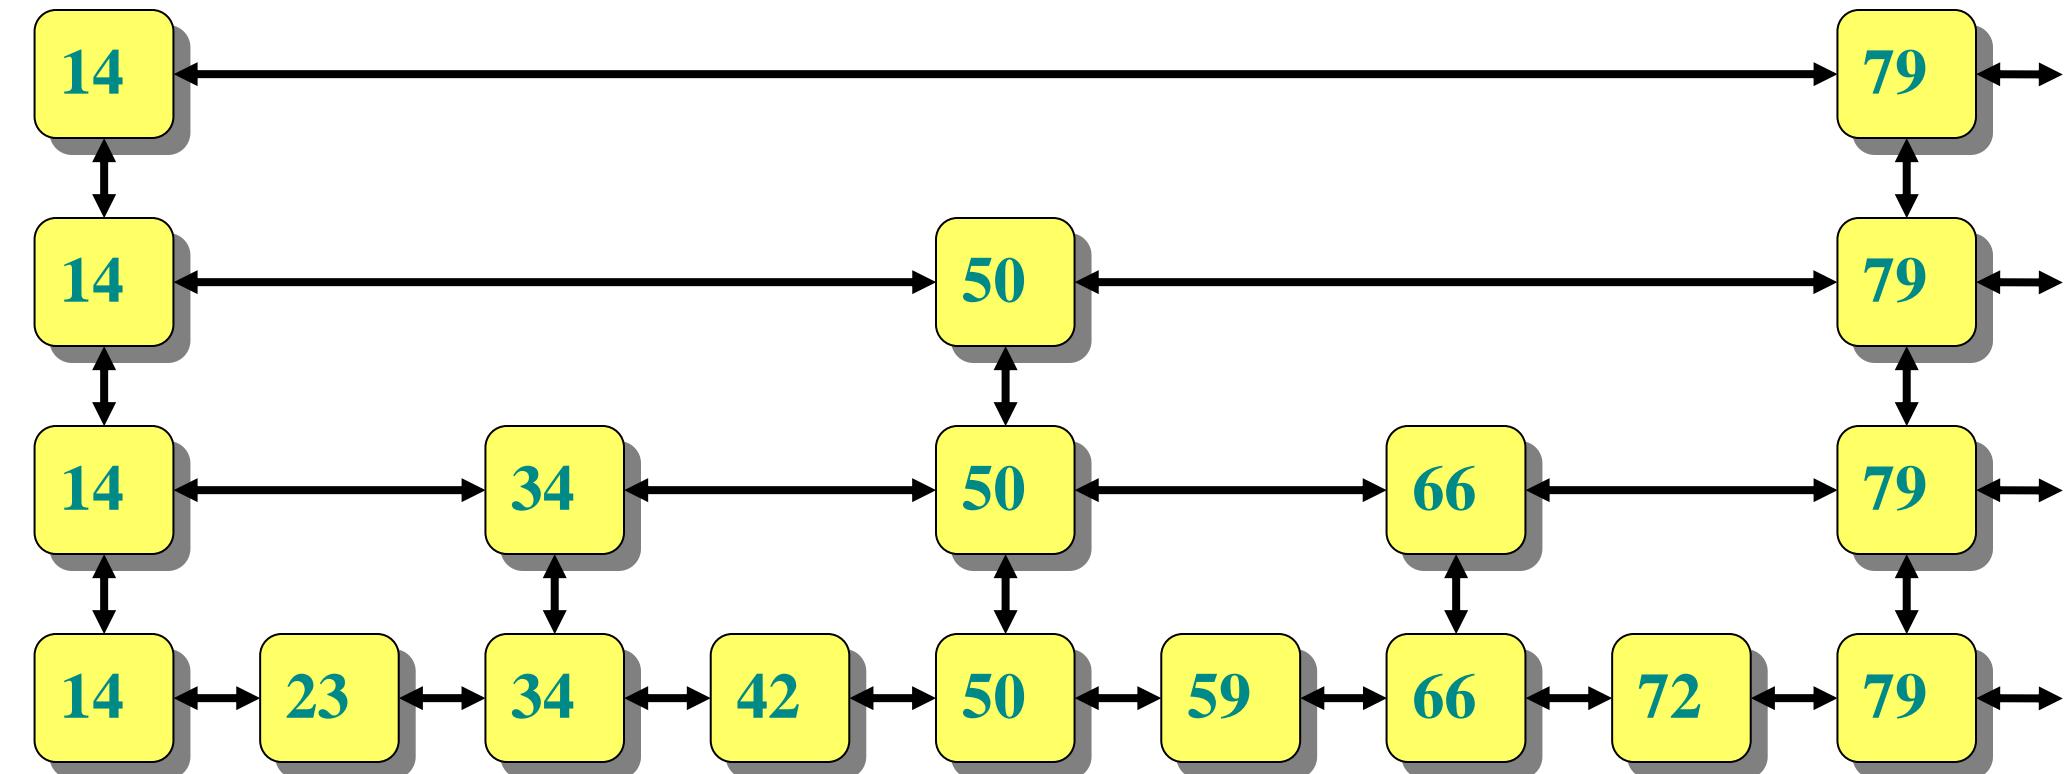
- 2 sorted lists $\Rightarrow 2 \cdot \sqrt{n}$
- 3 sorted lists $\Rightarrow 3 \cdot \sqrt[3]{n}$
- k sorted lists $\Rightarrow k \cdot \sqrt[k]{n}$
- $\lg n$ sorted lists $\Rightarrow \lg n \cdot \sqrt[\lg n]{n} = 2 \lg n$





$\lg n$ linked lists

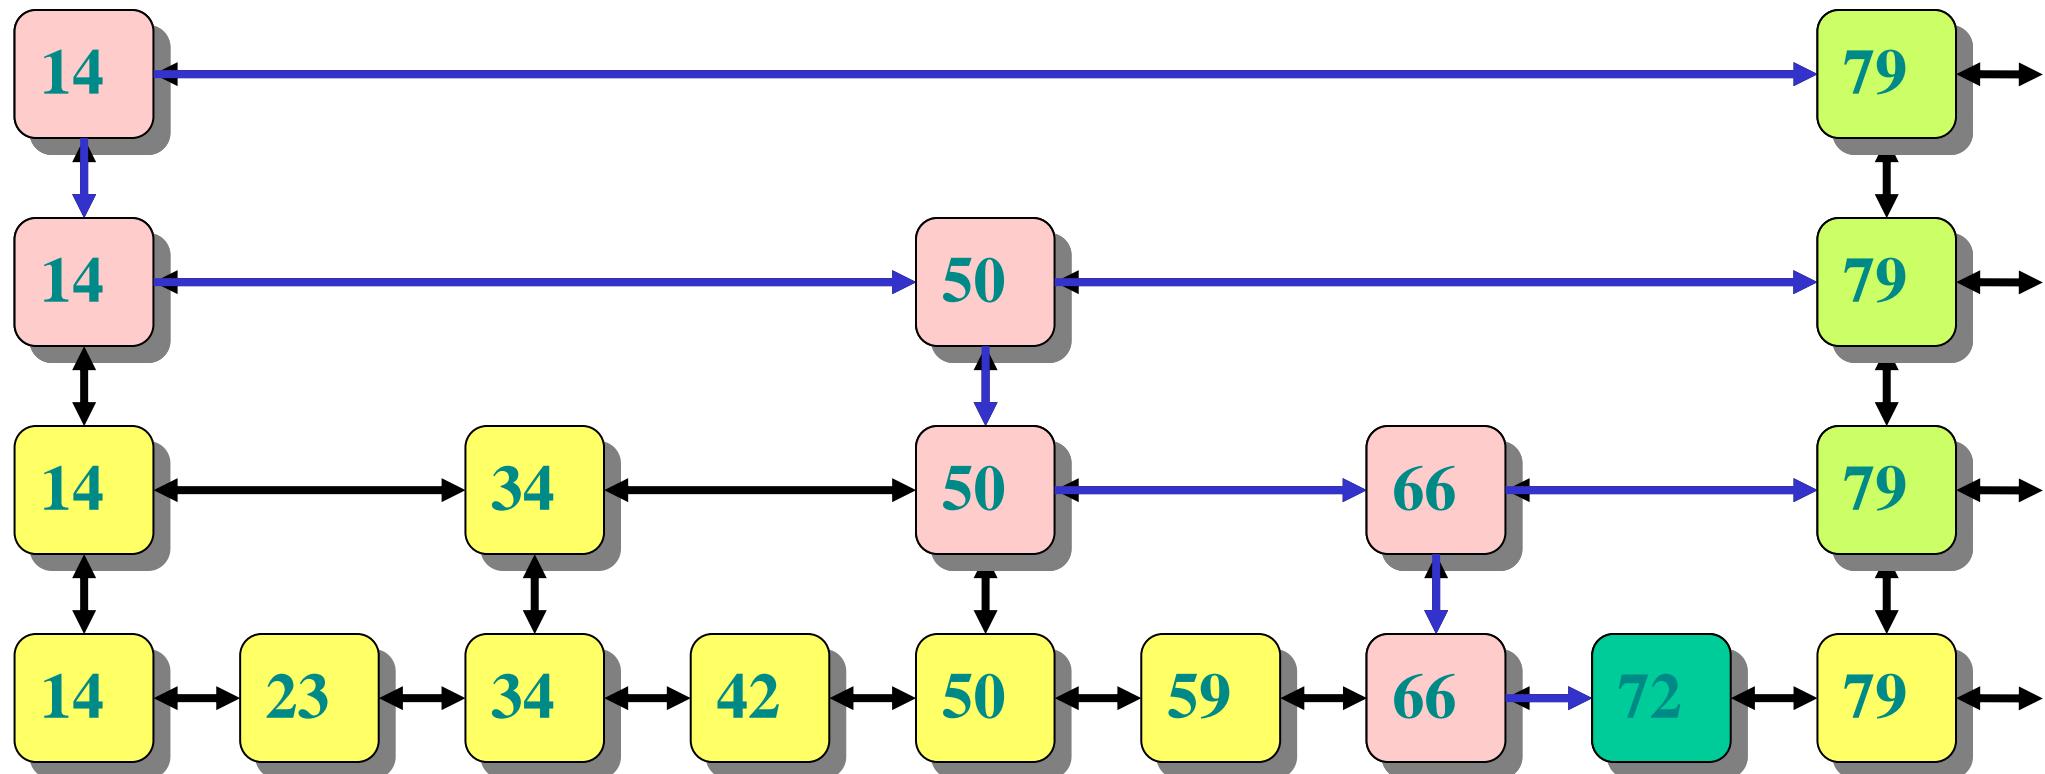
$\lg n$ sorted linked lists are like a binary tree





Searching in $\lg n$ linked lists

EXAMPLE: SEARCH(72)

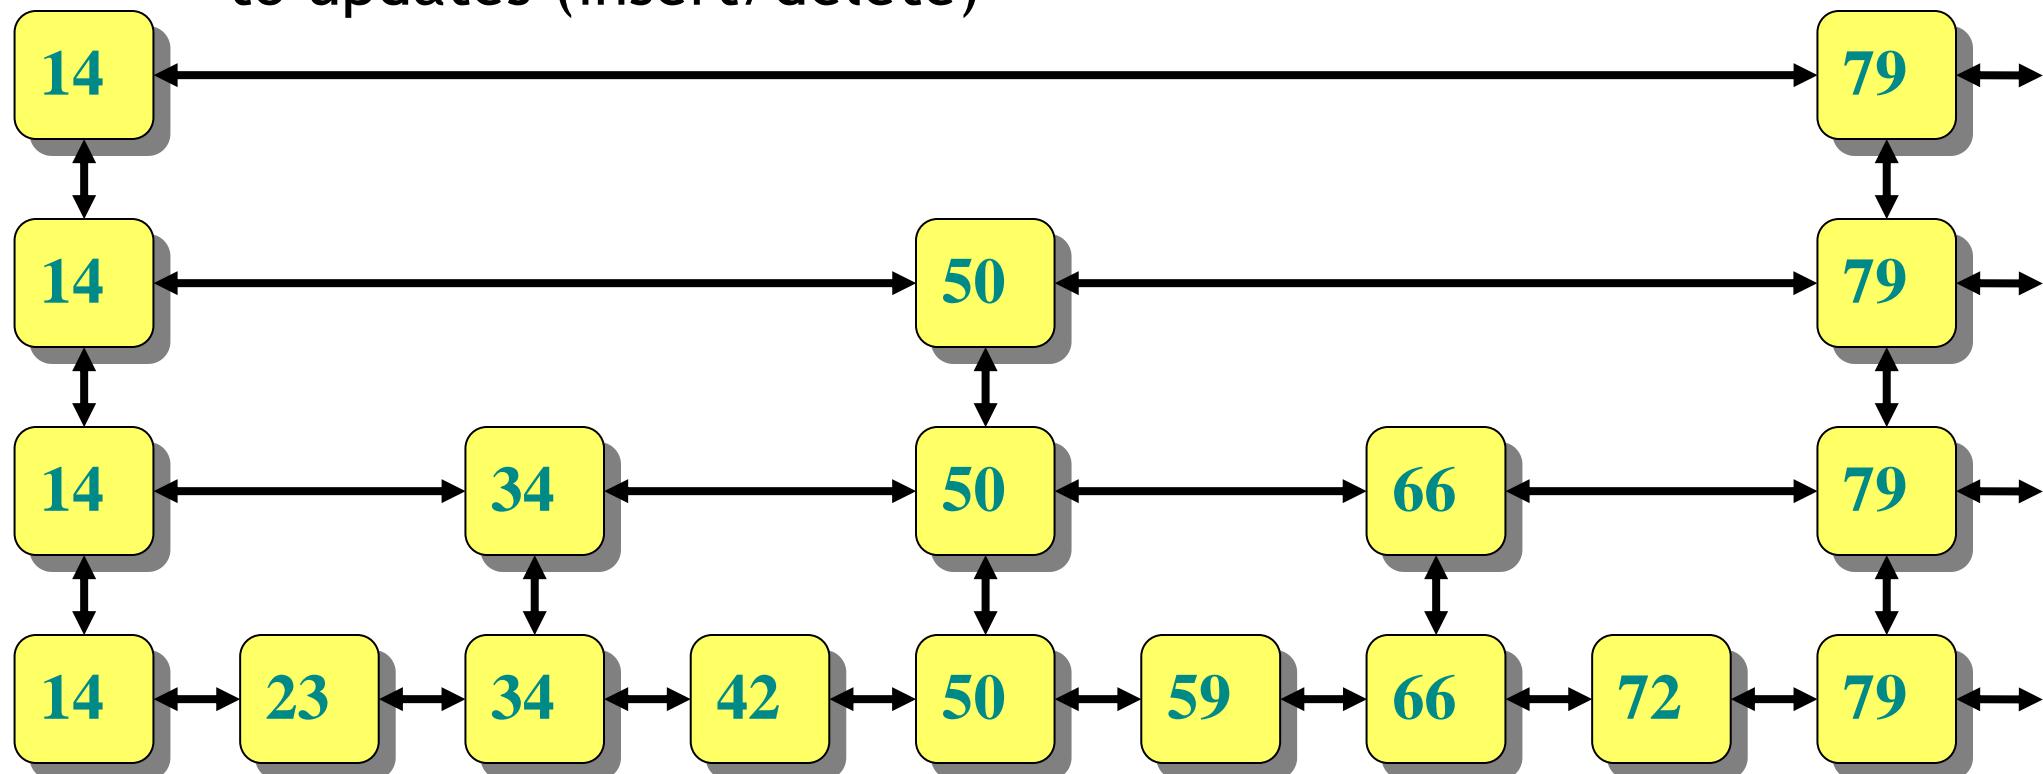




Skip lists

Ideal skip list is this $\lg n$ linked list structure

Skip list data structure maintains roughly this structure subject to updates (insert/delete)





INSERT(x)

To insert an element x into a skip list:

- ❑ SEARCH(x) to see where x fits in bottom list
- ❑ Always insert into bottom list

INVARIANT: Bottom list contains all elements

- ❑ Insert into some of the lists above...

QUESTION: To which other lists should we add x ?

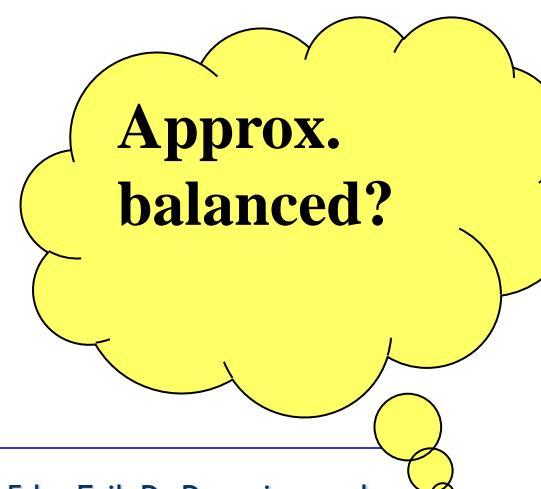


INSERT(x)

QUESTION: To which other lists should we add x ?

IDEA: Flip a (fair) coin; if HEADS,
promote x to next level up and flip again

- ❑ Probability of promotion to next level = 1/2
- ❑ On average:
 - 1/2 of the elements promoted 0 levels
 - 1/4 of the elements promoted 1 level
 - 1/8 of the elements promoted 2 levels
 - etc.



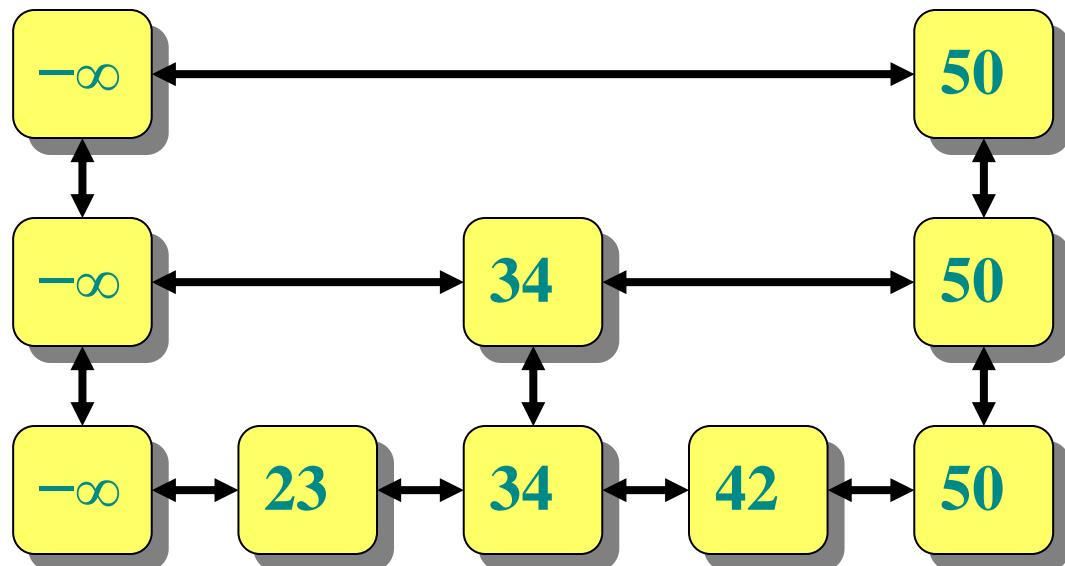


Example of skip list

EXERCISE: Try building a skip list from scratch
by repeated insertion using a real coin

Small change:

- Add special $-\infty$ value to *every* list
 \Rightarrow can search with the same algorithm





Skip lists

A *skip list* is the result of insertions (and deletions) from an initially empty structure (containing just $-\infty$)

- **INSERT(x)** uses random coin flips to decide promotion level
- **DELETE(x)** removes x from all lists containing it



Skip lists

A *skip list* is the result of insertions (and deletions) from an initially empty structure (containing just $-\infty$)

- $\text{INSERT}(x)$ uses random coin flips to decide promotion level
- $\text{DELETE}(x)$ removes x from all lists containing it

How good are skip lists? (speed/balance)

- **INTUITIVELY:** Pretty good on average
- **CLAIM:** Really, really good, almost always



With-high-probability theorem

THEOREM: *With high probability*, every search in an n -element skip list costs $O(\lg n)$



With-high-probability theorem

THEOREM: *With high probability*, every search in a skip list costs $O(\lg n)$

- **INFORMALLY:** Event E occurs ***with high probability*** (w.h.p.) if, for any $\alpha \geq 1$, there is an appropriate choice of constants for which
 E occurs with probability at least $1 - O(1/n^\alpha)$
 - In fact, constant in $O(\lg n)$ depends on α
- **FORMALLY:** Parameterized event E_α occurs ***with high probability*** if, for any $\alpha \geq 1$, there is an appropriate choice of constants for which
 E_α occurs with probability at least $1 - c_\alpha/n^\alpha$



With-high-probability theorem

THEOREM: With high probability, every search in a skip list costs $O(\lg n)$

- **INFORMALLY:** Event E occurs ***with high probability (w.h.p.)*** if, for any $\alpha \geq 1$, there is an appropriate choice of constants for which
 E occurs with probability at least $1 - O(1/n^\alpha)$
- **IDEA:** Can make ***error probability*** $O(1/n^\alpha)$ very small by setting α large, e.g., 100
- Almost certainly, bound remains true for entire execution of polynomial-time algorithm



Boole's inequality / union bound

Recall:

BOOLE'S INEQUALITY / UNION BOUND:

For any random events E_1, E_2, \dots, E_k ,

$$\begin{aligned} & \Pr\{E_1 \cup E_2 \cup \dots \cup E_k\} \\ & \leq \Pr\{E_1\} + \Pr\{E_2\} + \dots + \Pr\{E_k\} \end{aligned}$$

Application to with-high-probability events:

If $k = n^{O(1)}$, and each E_i occurs with high probability, then so does $E_1 \cap E_2 \cap \dots \cap E_k$



Analysis Warmup

LEMMA: With high probability,
 n -element skip list has $O(\lg n)$ levels

PROOF:

- Error probability for having at most $c \lg n$ levels
 - = $\Pr\{\text{more than } c \lg n \text{ levels}\}$
 - $\leq n \cdot \Pr\{\text{element } x \text{ promoted at least } c \lg n \text{ times}\}$
(by Boole's Inequality)
 - = $n \cdot (1/2^{c \lg n})$
 - = $n \cdot (1/n^c)$
 - = $1/n^{c-1}$



Analysis Warmup

LEMMA: With high probability,
 n -element skip list has $O(\lg n)$ levels

PROOF:

- Error probability for having at most $c \lg n$ levels
 $\leq 1/n^{c-1}$
- This probability is *polynomially small*,
i.e., at most n^α for $\alpha = c - 1$.
- We can make α arbitrarily large by choosing the constant c in
the $O(\lg n)$ bound accordingly.





Proof of theorem

THEOREM: With high probability, every search
in an n -element skip list costs $O(\lg n)$

COOL IDEA: Analyze search backwards—leaf to root

- Search starts [ends] at leaf (node in bottom level)
- At each node visited:
 - If node wasn't promoted higher (got TAILS here),
then we go [came from] left
 - If node was promoted higher (got HEADS here),
then we go [came from] up
- Search stops [starts] at the root (or $-\infty$)



Proof of theorem

THEOREM: With high probability, every search
in an n -element skip list costs $O(\lg n)$

COOL IDEA: Analyze search backwards—leaf to root

PROOF:

- Search makes “up” and “left” moves until it reaches the root (or $-\infty$)
- Number of “up” moves < number of levels
$$\leq c \lg n \text{ w.h.p. } (\text{Lemma})$$
- \Rightarrow w.h.p., number of moves is at most the number of times we need to flip a coin to get $c \lg n$ HEADS



Coin flipping analysis

CLAIM: Number of coin flips until $c \lg n$ HEADS
= $\Theta(\lg n)$ with high probability

PROOF:

Obviously $\Omega(\lg n)$: at least $c \lg n$

Prove $O(\lg n)$ “by example”:

- Say we make 10 $c \lg n$ flips
- When are there at least $c \lg n$ HEADS?

(Later generalize to arbitrary values of 10)



Coin flipping analysis

CLAIM: Number of coin flips until $c \lg n$ HEADS
= $\Theta(\lg n)$ with high probability

PROOF:

- $\Pr\{\text{exactly } c \lg n \text{ HEADS}\} =$
- $\Pr\{\text{at most } c \lg n \text{ HEADS}\} \leq$

$$\underbrace{\binom{10c \lg n}{c \lg n}}_{\text{orders}} \cdot \underbrace{\left(\frac{1}{2}\right)^{c \lg n}}_{\text{HEADS}} \cdot \underbrace{\left(\frac{1}{2}\right)^{9c \lg n}}_{\text{TAILS}}$$

$$\underbrace{\binom{10c \lg n}{c \lg n}}_{\text{overestimate on orders}} \cdot \underbrace{\left(\frac{1}{2}\right)^{9c \lg n}}_{\text{TAILS}}$$



Coin flipping analysis (cont'd)

- Recall bounds on

:

$$\binom{y}{x} \quad \left(\frac{y}{x}\right)^x \leq \binom{y}{x} \leq \left(e \frac{y}{x}\right)^x$$

- $\Pr\{\text{at most } c \lg n \text{ HEADS}\}$

$$\leq \binom{10c \lg n}{c \lg n} \cdot \left(\frac{1}{2}\right)^{9c \lg n}$$

$$\leq \left(e \frac{10c \lg n}{c \lg n}\right)^{c \lg n} \cdot \left(\frac{1}{2}\right)^{9c \lg n}$$

$$= (10e)^{c \lg n} 2^{-9c \lg n}$$

$$= 2^{\lg(10e) \cdot c \lg n} 2^{-9c \lg n}$$

$$= 2^{[\lg(10e) - 9] \cdot c \lg n}$$

$$= 1/n^\alpha \text{ for } \alpha = [9 - \lg(10e)] \cdot c$$



Coin flipping analysis (cont'd)

- $\Pr\{\text{at most } c \lg n \text{ HEADS}\} \leq 1/n^\alpha$ for $\alpha = [9 - \lg(10e)]c$
- **KEY PROPERTY:** $\alpha \rightarrow \infty$ as $10 \rightarrow \infty$, for any c
- So set 10, i.e., constant in $O(\lg n)$ bound,
large enough to meet desired α



This completes the proof of the coin-flipping claim
and the proof of the theorem.



Acknowledge

- Based on Randomized Algorithms, by R. Motwani and P. Raghavan.
- Based on Lectures of Prof. Dr. Th. Ottmann and of Prof. Dr. Susanne Albers: <http://lectures.informatik.uni-freiburg.de/portal/web/guest>
- Based on Lectures of Prof. Rada Mihalcea Univ. of North Texas